

Authority Engine Recovery Plan — From CMS Infrastructure to Agentic Content Factory

Executive Decision

We do not restart the project.

The current system is a strong foundation:

- Next.js + Payload CMS
- Authority graph
- Entity schema
- Lead capture
- Resend
- Player Trap funnel
- SEO/schema infrastructure
- Visibility baseline

But the original goal was broader:

Build a governed AI content production system that makes The Push and Itay Foyerstein discoverable, citable, and recommendable by AI answer engines.

The gap is clear: we built the publishing and authority infrastructure, but not the full 16-agent content factory originally envisioned.

The next phase must bridge that gap without breaking the verified production system.

North Star

Build a governed LangGraph-based agent factory that turns approved Itay insights into review-ready authority content, distribution assets, and measurable AI visibility improvements.

No agent may publish directly.

Payload remains the source of truth.

Human review remains mandatory before publication.

Current State

Already working:

- Marketing site
- Payload CMS
- Authority graph
- SEO/schema layer
- llms.txt / sitemap / robots.txt
- Player Trap funnel
- Lead capture
- Resend delivery
- UTM capture
- Visibility baseline
- Draft/review content states

Missing:

- 16 specialized agents
- LangGraph orchestration
- agent-specific typed contracts
- source verification
- framework builder
- brand voice enforcement
- distribution asset generation
- performance feedback loop
- recommendation pages
- original research assets
- AI visibility capture loop

Phase 0 — Freeze and Stabilize

Goal: Protect production while preparing the agent factory.

Actions: 1. Complete current in-progress hardening tasks. 2. Finish Task 023C without changing flow mechanics. 3. Complete Task 012A authority graph verification. 4. Complete Task 015 authority evidence layer. 5. Move Task 014 from review to completed after human approval. 6. Mark outdated Task 023B as superseded.

Definition of Done:

- Build passes.
- Typecheck passes.
- Existing Player Trap English/Hebrew flows still work.

- Authority graph verifier passes.
 - Release readiness checklist passes.
-

Phase 1 — Restore the Original Agent Architecture

Goal: Create the missing 16-agent architecture as code contracts before implementation.

Agents to define:

1. MarketIntelligenceAgent Finds trends, market shifts, competitor angles, emerging questions.
2. AudiencePainAgent Translates topics into real manager pain, internal voice, desired state, and emotional triggers.
3. SearchIntentAgent Classifies topics into awareness, consideration, decision, retention, and referral intent.
4. TopicClusterAgent Builds pillar/cluster architecture and assigns internal linking paths.
5. ResearchSourceAgent Finds sources, claims, definitions, statistics, and evidence.
6. SourceVerificationAgent Verifies that sources are real, current, relevant, and support the claims.
7. FrameworkBuilderAgent Turns ideas into proprietary frameworks such as Player Trap, Invisible Executor, Strategic Presence Map.
8. OutlineAgent Creates structured article outlines.
9. ContentWriterAgent Writes the draft in The Push voice.
10. LLMCitationAgent Adds short answers, definition boxes, citation snippets, target questions, FAQ, and quotable blocks.
11. InternalLinkingAgent Suggests internal links, anchor text, and entity connections.
12. SchemaAgent Generates schema recommendations.
13. BrandVoiceAgent Rejects generic coaching language and enforces sharp The Push positioning.
14. QualityGateAgent Scores usefulness, originality, support, structure, citation-worthiness, and conversion relevance.
15. PayloadPublisherAgent Saves only draft/review content into Payload.

16. DistributionAgent Creates LinkedIn post, carousel outline, email teaser, WhatsApp snippet, and short video script.
17. PerformanceLearningAgent Optional second-wave agent. Reads performance and recommends refreshes.

Architecture Decision: PerformanceLearningAgent may be implemented as Phase 2 if we want to keep the first bridge smaller.

Deliverables:

- `/src/ai/agents/*`
- `/src/ai/schemas/*`
- `/src/ai/prompts/*`
- `/src/ai/workflows/content-factory.workflow.ts`
- `/docs/AGENT_FACTORY.md`

Definition of Done:

- Every agent has typed input.
- Every agent has typed output.
- Every agent has a Zod schema.
- Every agent has failure states.
- No agent writes published content.
- Tests cover contract validation.

Phase 2 — LangGraph Orchestration

Goal: Build the orchestrator that manages the agents.

Canonical workflow:

Approved Itay Insight → Market Intelligence → Audience Pain → Search Intent → Topic Cluster → Research → Source Verification → Framework Builder → Outline → Content Writer → LLM Citation → Internal Linking → Schema → Brand Voice Review → Quality Gate → Payload Draft → Distribution Assets → Human Review

Hard rules:

- No approved insight = no content generation.
- Failed source verification = stop.
- Failed quality gate = return to previous agent.
- Failed brand voice = rewrite required.
- Failed human review = do not publish.
- PayloadPublisherAgent may only create draft/review records.

Deliverables:

- LangGraph workflow
- Agent run logging
- Prompt version tracking
- Retry/failure state handling
- Human review handoff

Definition of Done:

- One approved insight can produce one complete review-ready draft.
 - Draft is saved in Payload.
 - Agent runs are logged.
 - Distribution assets are attached or stored.
 - No automatic publishing occurs.
-

Phase 3 — Content Factory v1

Goal: Generate the first controlled batch of authority content.

Initial batch:

- 1 Pillar Page
- 12 Cluster Pages
- 20 FAQ entries
- 10 Glossary terms
- 5 Case Studies
- 3 Framework pages

Content rules:

- 70% educational
- 30% The Push positioning
- Mostly non-branded
- Every page answers 5–10 AI-style questions
- Every page includes short answer, key takeaways, definition box, tables, FAQ, schema, author credibility, last updated, internal links, and citation snippets.

Priority content: 1. The Player Trap 2. Invisible Executor 3. The Push Leadership Evolution Model 4. Tech Leadership Coaching 5. Engineering Manager Strategic Leadership 6. How to stop being the bottleneck 7. Tech Lead to Engineering Manager transition 8. Managing up as an engineering leader 9. Leadership visibility for technical managers 10. Why strong technical managers do not get promoted

Definition of Done:

- All generated content is review-ready.

- No fake claims.
 - No unsupported metrics.
 - All claims have evidence or are clearly framed as opinion/framework.
 - Human reviewer can approve, reject, or request changes.
-

Phase 4 — Recommendation Layer

Goal: Move from authority content to recommendation-intent pages.

Pages to create:

- /coach-for-tech-managers
- /engineering-manager-coaching
- /tech-lead-coach-israel
- /leadership-coach-for-rd-managers
- /why-itay-foyerstein
- /best-coach-for-engineering-managers
- /coach-for-managers-stuck-in-execution-mode

Each page must answer:

- Who is this for?
- When is Itay a good fit?
- When is Itay not a good fit?
- What problem does he solve?
- What proof exists?
- What framework does he use?
- What is the next step?

Definition of Done:

- Pages are not sales pages.
 - Pages are answer-engine friendly.
 - Pages support recommendation queries.
 - Pages connect to Player Trap and The Push.
 - Pages include clear CTA to diagnostic or fit call.
-

Phase 5 — Evidence and Proof Layer

Goal: Make the system credible enough for AI engines and humans.

Required assets: 1. Structured testimonials 2. Case studies 3. Before/after leadership stories 4. Evidence URLs 5. SameAs links 6. LinkedIn references 7. Client-type proof 8. Methodology proof 9. External mentions 10. Media/podcast/guest post records

Case study structure:

- Before
- Situation
- Intervention
- Result
- Timeframe
- Role
- Company type
- Evidence status
- Approved quote if available

Definition of Done:

- Every major claim has proof, source, or explicit evidence gap.
- Public pages expose evidence metadata.
- AI citation snippets are supported by visible evidence blocks.

Phase 6 — Distribution System

Goal: Every authority asset creates distribution assets automatically.

For every approved draft, generate:

- LinkedIn post
- LinkedIn carousel outline
- Email teaser
- WhatsApp post
- Short video script
- CTA variant
- Retargeting angle

Distribution does not auto-publish.

Human approves distribution copy.

Definition of Done:

- Every new content asset creates at least 3 distribution assets.
- Distribution assets link back to relevant page, Player Trap, or diagnostic call.
- LinkedIn becomes a systematic amplification channel, not manual improvisation.

Phase 7 — Measurement and Learning Loop

Goal: Close the loop between content, AI visibility, and leads.

Track:

- AI mentions
- AI citations
- query visibility
- page views
- CTA clicks
- lead source
- content source
- assessment completions
- diagnosis call requests
- booked calls
- close rate

PerformanceLearningAgent recommends:

- update title
- add FAQ
- add internal links
- create supporting cluster
- strengthen evidence
- improve CTA
- refresh content

Definition of Done:

- Every content asset has performance metadata.
- Every lead can be traced to source page / CTA / campaign where possible.
- Authority gaps produce content recommendations.
- No new content sprint starts without reviewing measurement data.

Governance Rules

1. Payload is the source of truth.
2. Agents may create drafts only.
3. Human review is mandatory for publishing.
4. No fresh approved Itay insight = no authority content generation.
5. No unsupported factual claims.
6. No invented testimonials.

7. No fake statistics.
 8. No generic coaching language.
 9. Every article must include at least one uncomfortable truth.
 10. Every task must have acceptance criteria before implementation.
-

Immediate Next Tasks

Task 025 — Agent Factory Contracts

Create typed contracts, Zod schemas, prompt shells, and documentation for the 16-agent content factory.

No runtime generation yet.

Task 026 — LangGraph Content Factory Orchestrator

Implement the governed workflow from approved insight to Payload draft.

Task 027 — Source Verification and Claim Ledger

Add claim/source verification before content writing.

Task 028 — First Agent-Generated Review Draft

Use one approved Itay insight to generate one review-ready article through the full pipeline.

Task 029 — Recommendation Pages Sprint

Create the first 5 recommendation-intent pages.

Task 030 — Distribution Asset Generator

Generate LinkedIn/email/WhatsApp/video assets from approved authority content.

Task 031 — Performance Learning Loop

Connect visibility and lead data back into content recommendations.

CTO Priority

Do not expand the Player Trap funnel further until the current copy hardening is complete.

Do not add the low-ticket book yet.

Do not build more CMS infrastructure unless required by the agent factory.

The next strategic bottleneck is:

Agentic content production with human governance.

That is the missing bridge between the system we built and the system we originally intended.